

Dummy GenAI Knowledge Document

Generative AI is a branch of artificial intelligence that can create new content such as text, code, images, and summaries. Large Language Models (LLMs) generate text by predicting likely tokens from the context provided to them.

Retrieval-Augmented Generation (RAG) improves an LLM application by retrieving relevant information from external documents before generating an answer. A typical RAG pipeline loads documents, splits them into chunks, creates embeddings, stores the embeddings in a vector store, retrieves relevant chunks, and sends the retrieved context to the language model.

Prompt templates make LLM applications easier to maintain. Instead of hard-coding every question, a template can contain placeholders such as {question}. The application fills those placeholders dynamically at runtime.

Chat prompt templates are useful for conversational applications because they separate system instructions, human messages, and optional AI messages. This structure makes the intended conversation flow explicit.

Example rule: Answer questions only from the supplied context. If the answer is not present in the context, say that the information is not available in the provided document.